

Ke Zhai

Machine Learning · Natural Language Processing · Cloud Computing
+1 (240)- 460-9082 · zhaikedavy@gmail.com · kzhai.github.io

EDUCATION

Ph.D. in Computer Science

University of Maryland, College Park, MD, 2014

GPA: 3.9/4.0. Supervisor: Jordan Boyd-Graber

Research topic: Large Scale Inference for Probabilistic Bayesian Models

Ph.D. thesis: Models, Inference, and Implementation for Scalable Probabilistic Models of Text

M.Sc in Computer Science

University of Maryland, College Park, MD, 2011

GPA: 3.9/4.0. Supervisor: Jordan Boyd-Graber and Jimmy Lin (co-supervised)

Research topic: Variational Bayesian Inference of Latent Dirichlet Allocation in MapReduce

Master scholarly paper: Using Variational Inference and MapReduce to Scale Topic Modeling

B.Eng. w. First Class Honor in Computer Engineering

Nanyang Technological University,

Singapore, 2009

GPA: 4.65/5.0. Supervisor: Wee Keong Ng

Research topic: Privacy-Preserving Data Mining

Undergraduate thesis: An Embedded Caching Framework for Privacy-Preserving Data Mining

EMPLOYMENT

Senior Research Scientist

Microsoft, Sunnyvale, CA, Aug 2016 - Present

- Research on language model and domain adaptation for speech recognition system.
- Language model personalization using topic models, context matching and style filtering.
- Maintaining end-to-end data processing, model training and adaptation evaluation pipeline.
- Language model training and adaptation support for Microsoft Teams project (a prototype of the product was demo'ed at Microsoft Build conference 2018).

Research Scientist

Yahoo! Labs, Sunnyvale, CA, Feb 2015 - Aug 2016

- Research on query understanding and sequence tagging for online Ads serving system.
- Research on query classification and user intent prediction for mobile search product.
- Research, design and develop on "Chat Bot as a Service" platform.

Graduate Research Assistant Department of Computer Science, University of Maryland, College Park, MD, Sep 2010 - Jun 2014

- Academic Advisor: Jordan Boyd-Graber
- Research Interest: Machine Learning, Non-parametric Bayesian Learning, Cloud Computing
- Design and implement online variational inference for adaptor grammars.
- Design and implement online variational inference for topic models with infinite vocabulary.
- Design and implement variational inference for latent Dirichlet allocation in MapReduce.
- Design and implement variational inference for Indian buffet process in MapReduce.

Research Intern

Yahoo! Labs, New York City, NY, Jun 2014 - Aug 2014

- Large scale unsupervised nonparametric models for user behavior analysis.

Research Intern

Microsoft Research, Redmond, WA, May 2013 - Aug 2013

- Mentor: Jason D. Williams

- Design and implement three models in discovering latent structure in dialogues.
- Achieve comparably well results against many other models on real datasets.

Software Engineering Intern

comScore, Inc., Reston, VA, May 2010 - Aug 2010

- Implement data transfer and formatter block for new deployed Hadoop distributed file system.
- Research and develop cookie deletion and prediction system for comScore.

Graduate Teaching Assistant Department of Computer Science, University of Maryland, College Park, MD, Sep 2009 - May 2010

- CMSC 131: Object-Oriented Programming I, Fall 2009.
- CMSC 422: Introduction to Artificial Intelligence, Spring 2010.

Software Engineering Industrial Attachment

StarHub, Inc., Singapore, Jan 2008 - Jun 2008

- Study on customer service system and client feedback process.
- Develop data collection and analysing software for Customer Service.
- Develop high-efficiency data processing system.
- Database development and administration.
- Design and develop reliable database system in MySQL for Customer Service.
- Automate data processing and analysing system with Microsoft Office and VBA.

PUBLICATION (* indicates equal contributor)

Ke Zhai*, and Huan Wang*. “Adaptive Dropout with Rademacher Complexity Regularization”. *International Conference on Learning Representations (ICLR)*, May 2018.

Ke Zhai, Zornitsa Kozareva, Yuening Hu, Qi Li and Weiwei Guo. “Query to Knowledge: Unsupervised Entity Extraction from Shopping Queries using Adaptor Grammars”. *International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR)*, Jul 2016.

Zornitsa Kozareva, Qi Li, **Ke Zhai** and Weiwei Guo. “Recognizing Salient Entities in Shopping Queries”. *Annual Meeting of the Association for Computational Linguistics (ACL)*, Jun 2016.

Ke Zhai, Jordan Boyd-Graber and Shay B. Cohen. “Online Adaptor Grammars with Hybrid Inference”. *Transaction of the Association for Computational Linguistics (TACL)*, Oct 2014.

Ke Zhai, and Jason D. Williams. “Discovering Latent Structure in Task-Oriented Dialogues”. *Annual Meeting of the Association for Computational Linguistics (ACL)*, Jun 2014.

Ke Zhai*, Yuening Hu*, Vladimir Edelman, and Jordan Boyd-Graber. “Polylingual Tree-Based Topic Models for Translation Domain Adaptation”. *Annual Meeting of the Association for Computational Linguistics (ACL)*, Jun 2014.

Ke Zhai, and Jordan Boyd-Graber. “Online Latent Dirichlet Allocation with Infinite Vocabulary”. *International Conference on Machine Learning (ICML)*, Jun 2013.

Ke Zhai*, Yuening Hu*, Jordan Boyd-Graber, and Sinead Williamson. “Modeling Images using Transformed Indian Buffet Processes”. *International Conference on Machine Learning (ICML)*, Jun 2012.

Ke Zhai, Jordan Boyd-Graber, Nima Asadi, and Mohamad Alkhouja. “Mr. LDA: A Flexible Large Scale Topic Modeling Package using Variational Inference in MapReduce”. *ACM International Conference on World Wide Web (WWW)*, Apr 2012.

Ke Zhai, Wee Keong Ng, Andre Ricardo Herianto and Shuguo Han. “Speeding Up Secure Computations via Embedded Caching”. *Proceedings of SIAM International Conference on Data Mining* (SDM), Apr 2009.

PROFESSIONAL CONTRIBUTION

Activate contributor for MapReduce library Cloud⁹ and Hadoop toolkit Ivory.

Online released code-base on GitHub, <https://github.com/kzhai>.

- MapReduce latent Dirichlet allocation (Variational Bayesian inference, with extension to informed prior and polylingual LDA).

“I tried both and found that, despite *Mr. LDA*’s cringeworthy name, it’s much faster than *Mahout*’s implementation so decided to go with that one.”

— **Kris Jack**, BSc Hons, Ph.D., Chief Data Scientist, Mendeley

- Latent Dirichlet allocation (Gibbs sampling, variational Bayesian inference and online version).
- Non-parametric Bayesian models (Indian buffet process, hierarchical Dirichlet process and infinite Gaussian mixture model).